

CNN Architecture Documentation

PyTorch Implementation Guide

September 16, 2025

Contents

1	Introduction	3
2	Imports and Setup	3
2.1	Library Functions	3
3	Data Preprocessing	3
3.1	Transformation Steps	3
4	Dataset Loading	4
5	Data Loading	4
6	Class Detection	4
7	CNN Architecture	4
7.1	Architecture Flow	5
7.2	Layer Analysis	5
8	Training Configuration	6
8.1	Configuration Details	6
9	Training Loop	6
9.1	Training Steps	6
10	Bug Fixes and Improvements	7
10.1	Identified Issues	7
10.2	Recommended Solutions	7
11	Performance Optimization	7
11.1	Training Improvements	7
11.2	Architecture Enhancements	7
12	Conclusion	8

1 Introduction

This document provides a comprehensive step-by-step breakdown of a Convolutional Neural Network (CNN) implementation in PyTorch for fruit classification. We'll examine each component from data loading to model training.

2 Imports and Setup

We begin by importing the essential PyTorch libraries and setting up our computational environment.

Listing 1: Library Imports and Device Configuration

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from torch.utils.data import DataLoader
5 from torchvision import datasets, transforms
6
7 device = torch.device("cpu")
```

2.1 Library Functions

- `torch` – Core PyTorch library for tensor operations
- `nn` – Neural network building blocks (Conv2d, Linear, etc.)
- `F` – Functional operations (ReLU, softmax, etc.)
- `DataLoader` – Efficient batch loading and shuffling
- `datasets` – Pre-built dataset classes like ImageFolder
- `transforms` – Image preprocessing and augmentation tools

Note: The device is set to CPU. For GPU acceleration, use:
`device = torch.device("cuda" if torch.cuda.is_available() else "cpu")`

3 Data Preprocessing

Image preprocessing is crucial for consistent model input and better convergence.

Listing 2: Image Transformation Pipeline

```
1 transform = transforms.Compose([
2     transforms.Resize((64,64)),
3     transforms.ToTensor(),
4     transforms.Normalize(mean=[0.5,0.5,0.5], std=[0.5,0.5,0.5])
5 ])
```

3.1 Transformation Steps

1. **Resize(64,64):** Standardizes all images to 64×64 pixels
2. **ToTensor():** Converts PIL images to PyTorch tensors (values 0-1)

3. **Normalize()**: Transforms values from [0,1] to [-1,1] using:

$$\text{normalized} = \frac{\text{pixel} - 0.5}{0.5} \quad (1)$$

4 Dataset Loading

We use PyTorch's ImageFolder class to automatically organize our fruit dataset.

Listing 3: Dataset Configuration

```
1 train_dir = '~/datasets/fruits/fruits-360_100x100/fruits-360/Training'
2 test_dir  = '~/datasets/fruits/fruits-360_100x100/fruits-360/Test'
3
4 train_dataset = datasets.ImageFolder(train_dir, transform = transform)
5 test_dataset  = datasets.ImageFolder(test_dir, transform = transform)
```

ImageFolder expects images organized in class-specific subfolders and automatically assigns labels based on folder names.

5 Data Loading

DataLoaders handle batch processing and data shuffling efficiently.

Listing 4: Batch Processing Setup

```
1 train_loader = DataLoader(train_dataset , batch_size = 64 ,shuffle=True)
2 test_loader  = DataLoader(test_dataset , batch_size = 64 ,shuffle=True)
```

Parameters:

- **batch_size = 64:** Process 64 images per iteration
- **shuffle=True:** Randomize order for better training

6 Class Detection

Listing 5: Automatic Class Counting

```
1 num_classes = len(train_dataset.classes)
2 print(f"Number of classes :{num_classes}")
```

This automatically counts the number of fruit categories from the folder structure.

7 CNN Architecture

The SimpleCNN class defines our neural network architecture.

Listing 6: CNN Model Definition

```
1 class SimpleCNN(nn.Module):
2     def __init__(self,num_classes):
3         super(SimpleCNN,self).__init__()
4         self.conv1 = nn.Conv2d(3,16,kernel_size=3,padding=1)    # input:
                        # 3 channels, output:16
5         self.conv2 = nn.Conv2d(16,32,kernel_size=3,padding=1)  # input
                        # :16 channels, output:32
```

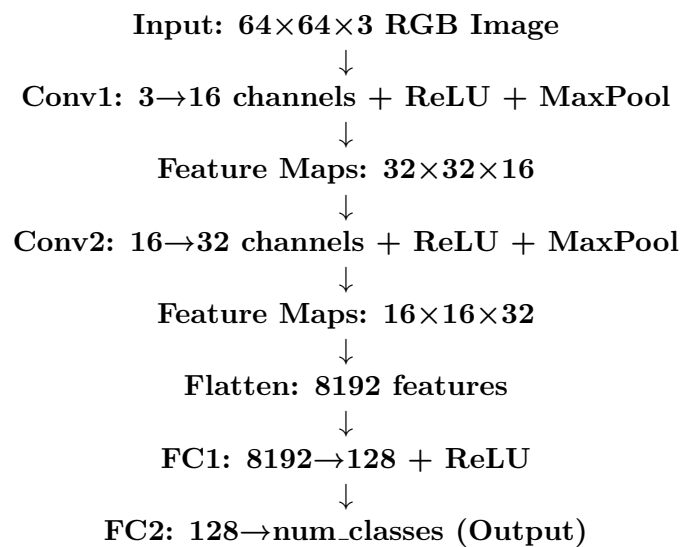
```

6         self.pool = nn.MaxPool2d(2,2)                                # reduces
           H,W by 2
7         self.fc1 = nn.Linear(32*16*16,128)                          # flatten
           to 32*16*16
8         self.fc2 = nn.Linear(128,num_classes)                       # final
           output layer
9
10        def forward(self,X):
11            X = self.pool(F.relu(self.conv1(X)))                    # conv1 -> ReLU -> pool
12            X = self.pool(F.relu(self.conv2(X)))                    # conv2 -> ReLU -> pool
13            X = X.view(X.size(0),-1)                                # flatten for fully
           connected layer
14            X = F.relu(self.fc1(X))                                  # hidden fully connected
           layer
15            X = self.fc2(X)                                          # output layer (logits)
16            return X

```

7.1 Architecture Flow

The network processes data through these sequential stages:



7.2 Layer Analysis

Layer	Input Shape	Output Shape	Parameters
Input	—	64×64×3	0
Conv1	64×64×3	64×64×16	448
MaxPool1	64×64×16	32×32×16	0
Conv2	32×32×16	32×32×32	4,640
MaxPool2	32×32×32	16×16×32	0
FC1	8,192	128	1,048,704
FC2	128	num_classes	128 × classes

Table 1: Layer-wise Architecture Breakdown

Key Points:

- Convolutional layers extract spatial features

- MaxPooling reduces dimensions while preserving important information
- Fully connected layers perform final classification
- Total flattened size: $32 \times 16 \times 16 = 8,192$ features

8 Training Configuration

Listing 7: Model and Training Setup

```
1 model = SimpleCNN(num_classes).to(device)
2 criterion = nn.CrossEntropyLoss()
3 optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
4 num_epochs = 5
```

8.1 Configuration Details

- **CrossEntropyLoss:** Standard loss for multi-class classification
- **Adam Optimizer:** Adaptive learning with momentum
- **Learning Rate = 0.001:** Conservative but stable rate
- **Epochs = 5:** Quick training for initial testing

9 Training Loop

Listing 8: Model Training Process

```
1 for epoch in range(num_epochs):
2     model.train()
3     running_loss = 0.0
4     for images, labels in train_loader:
5         images, labels = images.to(device), labels.to(device)
6         optimizer.zero_grad()           # clear previous gradients
7         outputs = model(images)          # forward pass
8         loss = criterion(outputs, labels) # compute loss
9         loss.backward()                  # backpropagation
10        optimizer.step()                  # update weights
11        running_loss += loss.item()
12    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(
        train_loader):.4f}")
```

9.1 Training Steps

The training follows this standard deep learning workflow:

1. **Zero Gradients:** Clear gradients from previous iteration
2. **Forward Pass:** Compute model predictions
3. **Loss Calculation:** Compare predictions with true labels
4. **Backward Pass:** Compute gradients via backpropagation
5. **Parameter Update:** Apply optimizer to update weights
6. **Loss Tracking:** Monitor training progress

10 Bug Fixes and Improvements

10.1 Identified Issues

1. **Typo:** `num_eporchs` should be `num_epochs`
2. **Path Issue:** Tilde (~) may not expand in all environments
3. **Missing Validation:** No validation loop to monitor overfitting

10.2 Recommended Solutions

Fix Path Expansion:

```
1 import os
2 train_dir = os.path.expanduser('~/.datasets/fruits/fruits-360_100x100/
   fruits-360/Training')
3 test_dir = os.path.expanduser('~/.datasets/fruits/fruits-360_100x100/
   fruits-360/Test')
```

Add Validation Loop:

```
1 # Add after training loop
2 model.eval()
3 correct = 0
4 total = 0
5 with torch.no_grad():
6     for images, labels in test_loader:
7         images, labels = images.to(device), labels.to(device)
8         outputs = model(images)
9         _, predicted = torch.max(outputs.data, 1)
10        total += labels.size(0)
11        correct += (predicted == labels).sum().item()
12
13 print(f'Test Accuracy: {100 * correct / total:.2f}%')
```

11 Performance Optimization

11.1 Training Improvements

- Implement learning rate scheduling
- Add early stopping mechanism
- Use data augmentation (rotation, flipping, cropping)
- Add batch normalization for stable training
- Implement dropout for regularization

11.2 Architecture Enhancements

- Add more convolutional layers for deeper feature extraction
- Experiment with different kernel sizes
- Try residual connections (ResNet-style)
- Use pre-trained models for transfer learning
- Implement attention mechanisms

12 Conclusion

This CNN implementation provides a solid foundation for image classification tasks. The architecture effectively combines convolutional feature extraction with fully connected classification layers.

Key Takeaways:

- Proper data preprocessing ensures consistent model input
- Convolutional layers learn hierarchical visual features
- Pooling operations provide translation invariance and reduce computation
- The training loop follows standard gradient descent optimization
- Regular validation prevents overfitting

The documented code is ready for implementation and can serve as a baseline for more advanced architectures and techniques.

Created by **Mohamed Amine**

GitHub: <https://github.com/Amine-DevAI>

LinkedIn: <https://www.linkedin.com/in/mohamed-amine-mammar-el-hadj-715a41295>

September 16, 2025